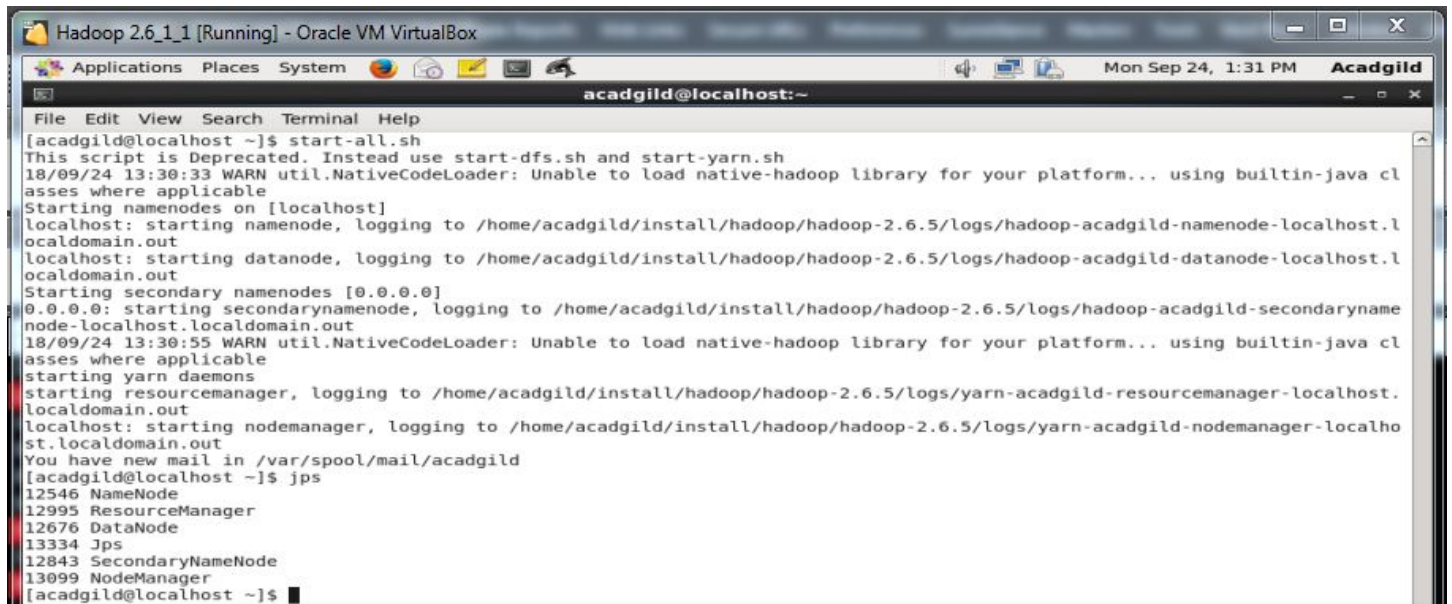


Session 9: Advance Hive Assignment 1.

1. To perform the Hive operations, first we need to start the hadoop in VM as follows:

\$start-all.sh



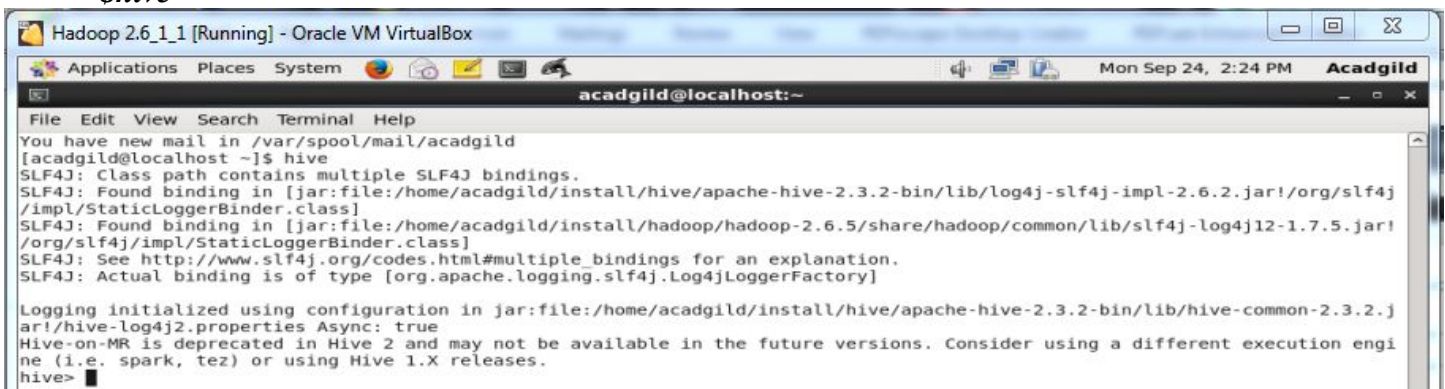
```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
acadgild@localhost:~
File Edit View Search Terminal Help
[acadgild@localhost ~]$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
18/09/24 13:30:33 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/hadoop-acadgild-namenode-localhost.l
ocaldomain.out
localhost: starting datanode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/hadoop-acadgild-datanode-localhost.l
ocaldomain.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/hadoop-acadgild-secondaryname
node-localhost.localdomain.out
18/09/24 13:30:55 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
starting yarn daemons
starting resourcemanager, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/yarn-acadgild-resourcemanager-localhost
.localdomain.out
localhost: starting nodemanager, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/yarn-acadgild-nodemanager-localhost
.localdomain.out
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ jps
12546 NameNode
12995 ResourceManager
12676 DataNode
13334 Jps
12843 SecondaryNameNode
13099 NodeManager
[acadgild@localhost ~]$
```

2. Check wither all the daemons are started in the hadoop as shown in the figure with the command:

\$jps

3. Now start the hive terminal with the following command:

\$hive

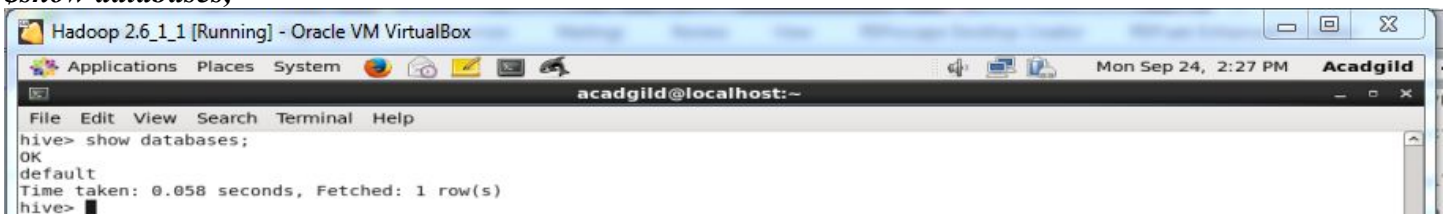


```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
acadgild@localhost:~
File Edit View Search Terminal Help
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j
/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!
/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/hive-common-2.3.2.j
ar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engi
ne (i.e. spark, tez) or using Hive 1.X releases.
hive>
```

4. Now we can see what are all the databases available, with the following command:

\$show databases;

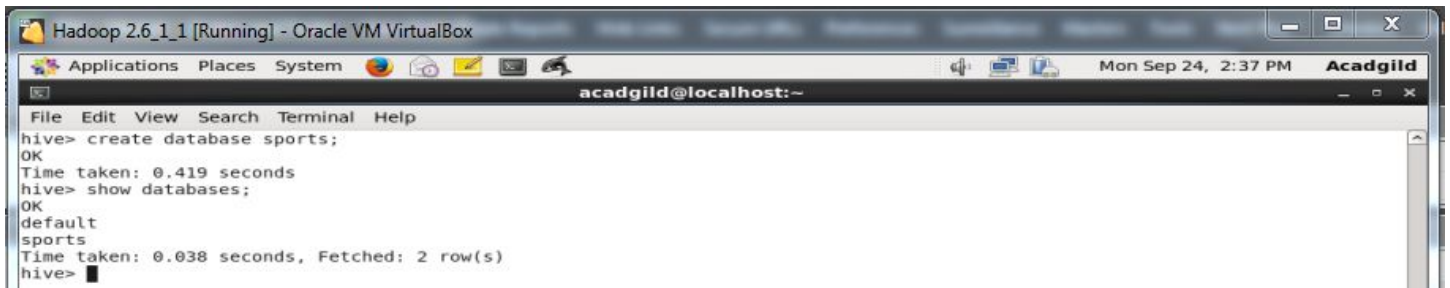


```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
acadgild@localhost:~
File Edit View Search Terminal Help
hive> show databases;
OK
default
Time taken: 0.058 seconds, Fetched: 1 row(s)
hive>
```

5. We can observe that there is only database available, i.e., default.

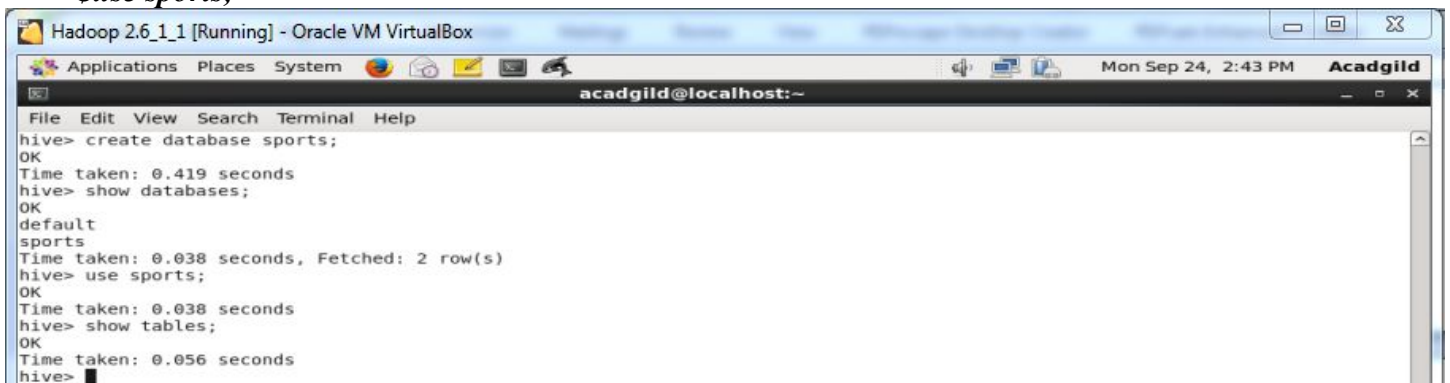
6. Now we can create a new database called **sports** as follows:

hive>create database sports;



```
Hadoop 2.6_1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadgild@localhost:~
File Edit View Search Terminal Help
hive> create database sports;
OK
Time taken: 0.419 seconds
hive> show databases;
OK
default
sports
Time taken: 0.038 seconds, Fetched: 2 row(s)
hive>
```

7. No to create tables we need to select the database lets say *sports*:
Use sports;




```
Hadoop 2.6_1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadgild@localhost:~
File Edit View Search Terminal Help
hive> create database sports;
OK
Time taken: 0.419 seconds
hive> show databases;
OK
default
sports
Time taken: 0.038 seconds, Fetched: 2 row(s)
hive> use sports;
OK
Time taken: 0.038 seconds
hive> show tables;
OK
Time taken: 0.056 seconds
hive>
```

8. We can also see that there are no tables present in the database *sports*.
9. Now we can create a table called *olympics* inside the database *sports* as follows:

CREATE TABLE OLYMPICS (

ath_name string,
ath_age int,
ath_country string,
year int,
closing_date string,
sport string,
gold_medals int,
silver_medals int,
bronze_medals int,
total_medals int

)Row format delimited fields terminated by '\t';



```
Hadoop 2.6_1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadgild@localhost:~
File Edit View Search Terminal Tabs Help
acadgild@localhost:~
hive> CREATE TABLE OLYMPICS (
>   ath_name string,
>   ath_age int,
>   ath_country string,
>   year int,
>   closing_date string,
>   sport string,
>   gold_medals int,
>   silver_medals int,
>   bronze_medals int,
>   total_medals int
> ) Row format delimited fields terminated by '\t';
OK
Time taken: 0.202 seconds
hive> show tables;
OK
olympics
Time taken: 0.063 seconds, Fetched: 1 row(s)
hive>
```

10. We can observe that the table is create with name Olympics and we can verify that table creation with *show tables* command as follows:

hive> show tables.

11. Now enter load data from the .csv file into the hive table as follows:

hive> load data local inpath '/home/acadgild/Downloads/olympix_data.csv' into table Olympics;

```
Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox
Applications Places System
*Unsaved Document 1 - gedit
File Edit View Search Tools Documents Help
+Unsaved Document 1
CREATE TABLE OLYMPICS (
  ath name string,
  ath age int,
  ath_country string,
  year int,
  closing_date string,
  sport string,
  gold_medals int,
  silver_medals int,
  bronze_medals int,
  total_medals int
) Row format delimited fields terminated by '\t';
load data local inpath '/home/acadgild/Downloads/olympix_data.csv' into table olympics;
```

12. Verify the records present in the table with *selec* command as follows:

*\$\$select * from olympics;*

```
Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox
Applications Places System
acadgild@localhost:~
File Edit View Search Terminal Tabs Help
acadgild@localhost:~
hive> load data local inpath '/home/acadgild/Downloads/olympix_data.csv' into table olympics;
Loading data to table sports.olympics
OK
Time taken: 0.923 seconds
hive> select * from olympics;
OK
Michael Phelps 23 United States 2008 08-24-08 Swimming 8 0 0 8
Michael Phelps 19 United States 2004 08-29-04 Swimming 6 0 2 8
Michael Phelps 27 United States 2012 08-12-12 Swimming 4 2 0 6
Natalie Coughlin 25 NULL 2008 NULL Swimming 1 3 6 NULL
Aleksey Nemov 24 Russia NULL 2000 10-01-00 NULL 2 1 3
Alicia Coutts 24 Australia 2012 08-12-12 Swimming 1 3 1 5
Missy Franklin 17 United States 2012 08-12-12 Swimming 4 0 1 5
Ryan Lochte 27 United States 2012 08-12-12 Swimming 2 2 1 5
Allison Schmitt 22 United States 2012 08-12-12 Swimming 3 1 1 5
Natalie Coughlin 21 NULL 2004 NULL Swimming 2 2 1 5
Ian Thorpe 17 Australia 2000 10-01-00 Swimming 3 2 0 5
Dara Torres 33 United States 2000 10-01-00 Swimming 2 0 3 5
Cindy Klassen 26 Canada NULL 2006 02-26-06 NULL 1 2 2
Nastia Liukin 18 United States 2008 08-24-08 Gymnastics 1 3 1 5
Marit Bjørgen 29 Norway NULL 2010 02-28-10 NULL 3 1 1
Sun Yang 20 China NULL 2012 08-12-12 NULL 2 1 1
Kirsty Coventry 24 Zimbabwe 2008 08-24-08 Swimming 1 3 0 4
Libby Lenton-Trickett 23 Australia 2008 08-24-08 Swimming 2 2 1 4
Ryan Lochte 24 United States 2008 08-24-08 Swimming 2 0 2 4
Inge de Bruijn 30 Netherlands 2004 08-29-04 Swimming 1 1 2 4
Petria Thomas 28 Australia 2004 08-29-04 Swimming 3 1 0 4
Ian Thorpe 21 Australia 2004 08-29-04 Swimming 2 1 1 4
Inge de Bruijn 27 Netherlands 2000 10-01-00 Swimming 3 1 0 4
Gary Hall Jr. 25 United States 2000 10-01-00 Swimming 2 1 1 4
Michael Klim 23 Australia 2000 10-01-00 Swimming 2 2 0 4
Susie O'Neill 27 Australia 2000 10-01-00 Swimming 1 3 0 4
Jenny Thompson 27 United States 2000 10-01-00 Swimming 3 0 1 4
Pieter van den Hoogenband 22 Netherlands 2000 10-01-00 Swimming 2 0 2 4
An Hyeon-Su 20 South Korea 2006 02-26-06 Short-Track Speed Skating 3 0 1 4
Aliya Mustafina 17 Russia 2012 08-12-12 Gymnastics 1 2 4 4
Shawn Johnson 16 United States 2008 08-24-08 Gymnastics 1 3 0 4
```

Task 1:

1. Write a Hive program to find the number of medals won by each country in swimming.

1. For this we can write a Hive query as follows:

hive> select ath_country, sum(total_medals) from Olympics where sport='Swimming' group by ath_country;


```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadgild@localhost:~
File Edit View Search Terminal Tabs Help
acadgild@localhost:~ acadgild@localhost:~
hive> select ath_country, sum(total_medals) from olympics where sport='Swimming' group by ath_country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180924171040_74012b6b-86ac-4af8-842e-294fa5b427b0
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537776060511_0004, Tracking URL = http://localhost:8088/proxy/application_1537776060511_0004/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537776060511_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-09-24 17:10:49,926 Stage-1 map = 0%, reduce = 0%
2018-09-24 17:10:58,748 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.34 sec
2018-09-24 17:11:08,613 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.91 sec
MapReduce Total cumulative CPU time: 7 seconds 910 msec
Ended Job = job_1537776060511_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.91 sec HDFS Read: 528840 HDFS Write: 881 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 910 msec
OK
Argentina 1
Australia 163
Austria 3
Belarus 2
Brazil 8
Canada 5
China 31
Costa Rica 2
Croatia 1
Denmark 1
France 39
Germany 32
Great Britain 11
Hungary 9
Italy 16
Japan 43
Lithuania 1
Netherlands 46
Norway 2
Poland 3
Romania 6
Russia 20
Serbia 1
Slovakia 2
Slovenia 1
South Africa 11
South Korea 4
Spain 3
Sweden 9
Trinidad and Tobago 1
Tunisia 3
Ukraine 7
United States 256
Zimbabwe 7
Time taken: 29.739 seconds, Fetched: 34 row(s)
hive>
```

2. We can observe the output that each country and corresponding total number of medals.
3. **Write a hive program to find the number of medals that India won year wise.**
 1. For this task we need to write a hive query as follows:
hive>select year, sum(total_medals) from Olympics where ath_country='India' group by year;

```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadmild@localhost:~
File Edit View Search Terminal Tabs Help
acadmild@localhost:~
hive> select year, sum(total_medals) from olympics where ath_country='India' group by year;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadmild_20180924171734_90107f48-2667-4f85-8f52-e835e9c4c2be
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537776060511_0005, Tracking URL = http://localhost:8088/proxy/application_1537776060511_0005/
Kill Command = /home/acadmild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537776060511_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-09-24 17:17:43,811 Stage-1 map = 0%, reduce = 0%
2018-09-24 17:17:52,644 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.69 sec
2018-09-24 17:18:01,508 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 8.45 sec
MapReduce Total cumulative CPU time: 8 seconds 450 msec
Ended Job = job_1537776060511_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.45 sec HDFS Read: 528829 HDFS Write: 163 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 450 msec
OK
2000      1
2004      1
2008      3
2012      6
Time taken: 29.333 seconds, Fetched: 4 row(s)
hive>
```

2. With the above figure we can see that the total number of medals own by India on each year.

3. Write a hive program to find the total number of medals each country own.

1. For this we need to write Hive query as follows:

hive> select ath_country, sum(total_medals) from Olympics group by ath_country;

```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadmild@localhost:~
File Edit View Search Terminal Tabs Help
acadmild@localhost:~
hive> select ath_country, sum(total_medals) from olympics group by ath_country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadmild_20180924172531_52d1b818-dbdb-4513-a0b5-51fef6a6c285
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537776060511_0006, Tracking URL = http://localhost:8088/proxy/application_1537776060511_0006/
Kill Command = /home/acadmild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537776060511_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-09-24 17:25:41,249 Stage-1 map = 0%, reduce = 0%
2018-09-24 17:25:48,924 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.92 sec
2018-09-24 17:25:58,689 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.01 sec
MapReduce Total cumulative CPU time: 6 seconds 10 msec
Ended Job = job_1537776060511_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.01 sec HDFS Read: 528008 HDFS Write: 2782 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 10 msec
OK
2004      NULL
2008      NULL
Afghanistan      2
Algeria      8
Argentina      141
Armenia      10
Australia      609
Austria      91
Azerbaijan      25
Bahamas      24
Bahrain      1
Barbados      1
```

```

Belarus 97
Belgium 18
Botswana 1
Brazil 221
Bulgaria 41
Cameroon 20
Canada 367
Chile 22
China 527
Chinese Taipei 20
Colombia 13
Costa Rica 2
Croatia 81
Cuba 188
Cyprus 1
Czech Republic 81
Denmark 89
Dominican Republic 5
Ecuador 1
Egypt 8
Eritrea 1
Estonia 18
Ethiopia 29
Finland 118
France 318
Gabon 1
Georgia 23
Germany 629
Great Britain 322
Greece 59
Grenada 1
Guatemala 1
Hong Kong 3
Hungary 145
Iceland 15
India 11
Indonesia 22
Indonesia 22
Iran 24
Ireland 9
Israel 4
Italy 331
Jamaica 80
Japan 282
Kazakhstan 42
Kenya 39
Kuwait 2
Kyrgyzstan 3
Latvia 17
Lithuania 30
Macedonia 1
Malaysia 3
Mauritius 1
Mexico 38
Moldova 5
Mongolia 10
Montenegro 14
Morocco 11
Mozambique 1
Netherlands 318
New Zealand 52
Nigeria 39
North Korea 21
Norway 188
Panama 1
Paraguay 17
Poland 80
Portugal 9
Puerto Rico 2
Qatar 3
Romania 123
Russia 765
Saudi Arabia 6
Serbia 31
Serbia and Montenegro 38
Singapore 7
Slovakia 35
Slovenia 25
South Africa 25
South Korea 308
Spain 205
Sri Lanka 1
Sudan 1
Sweden 181
Switzerland 93
Syria 1
Tajikistan 3
Thailand 18
Togo 1
Trinidad and Tobago 19
Tunisia 4
Turkey 28
Uganda 1
Ukraine 143
United Arab Emirates 1
United States 1301
Uruguay 1
Uzbekistan 19
Venezuela 4
Vietnam 2
Zimbabwe 7
Time taken: 28.643 seconds, Fetched: 112 row(s)
hive>

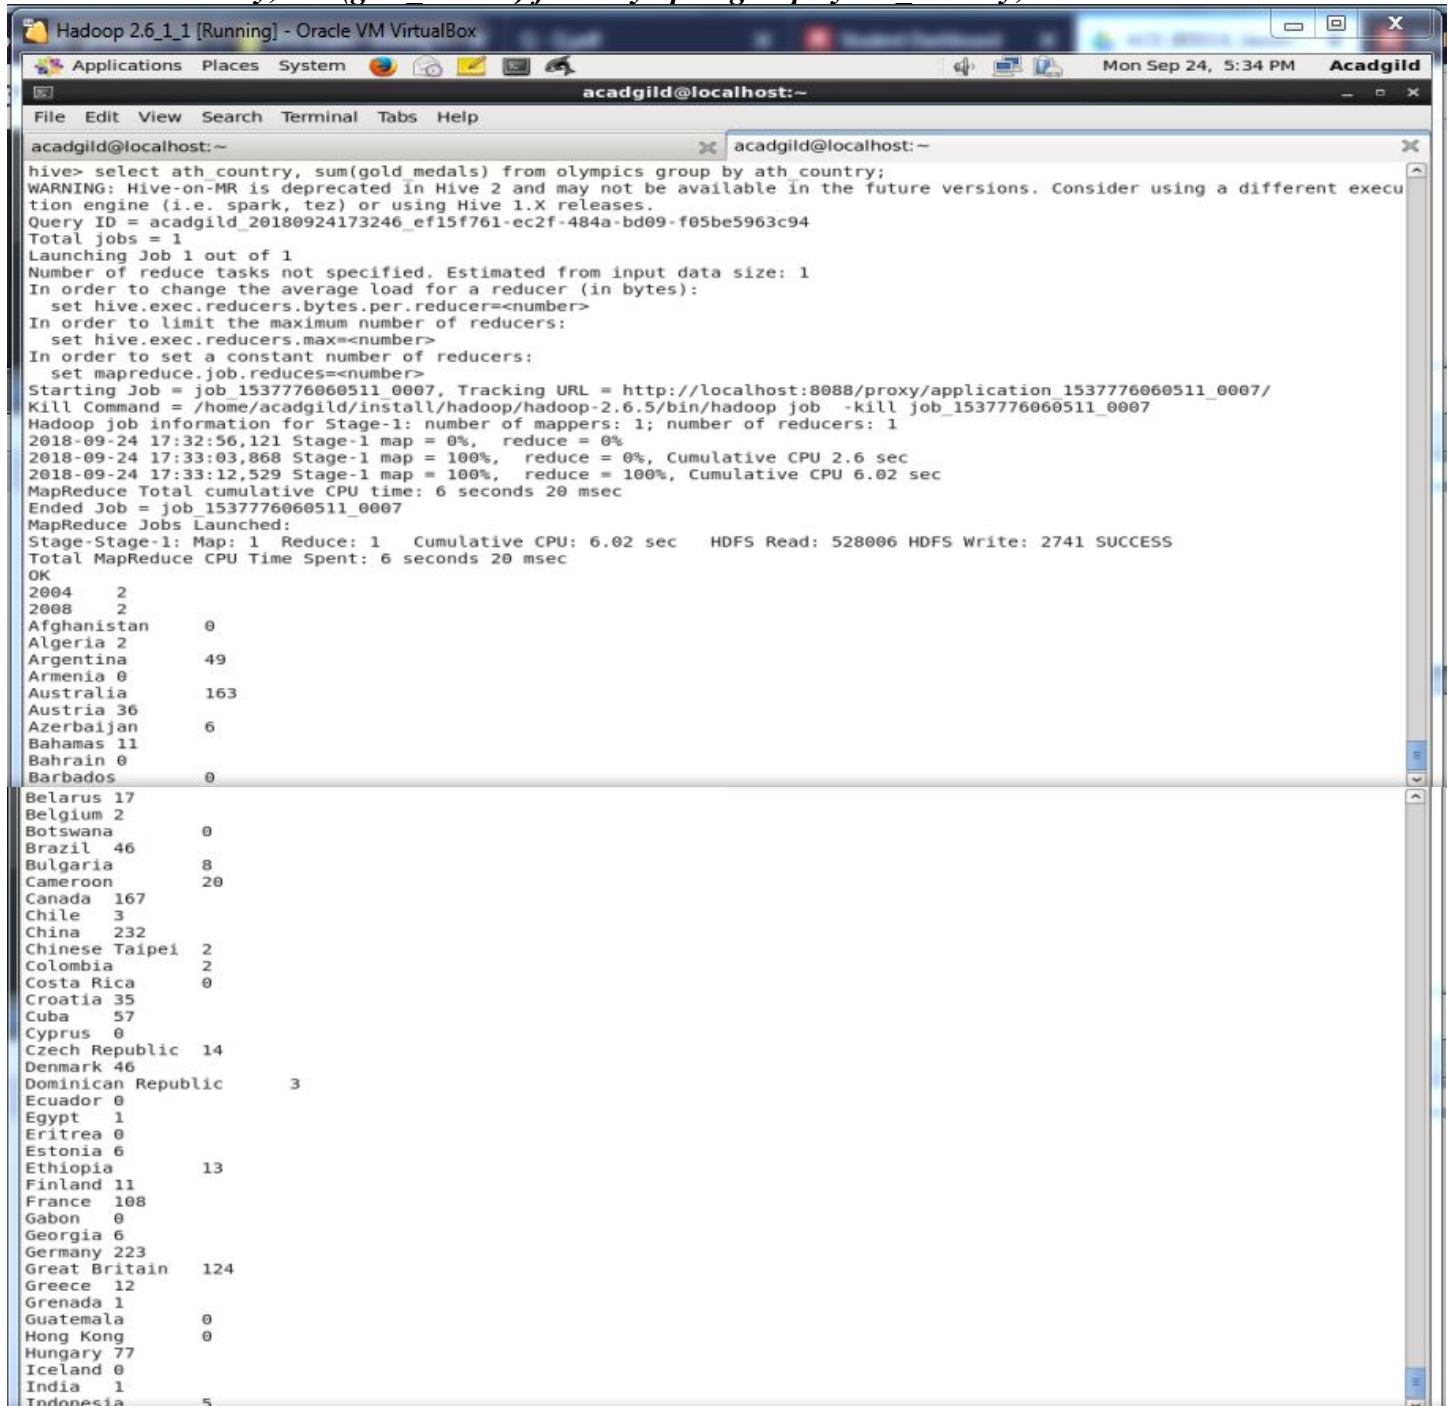
```

2. We can observe that each country and their number of medals own.

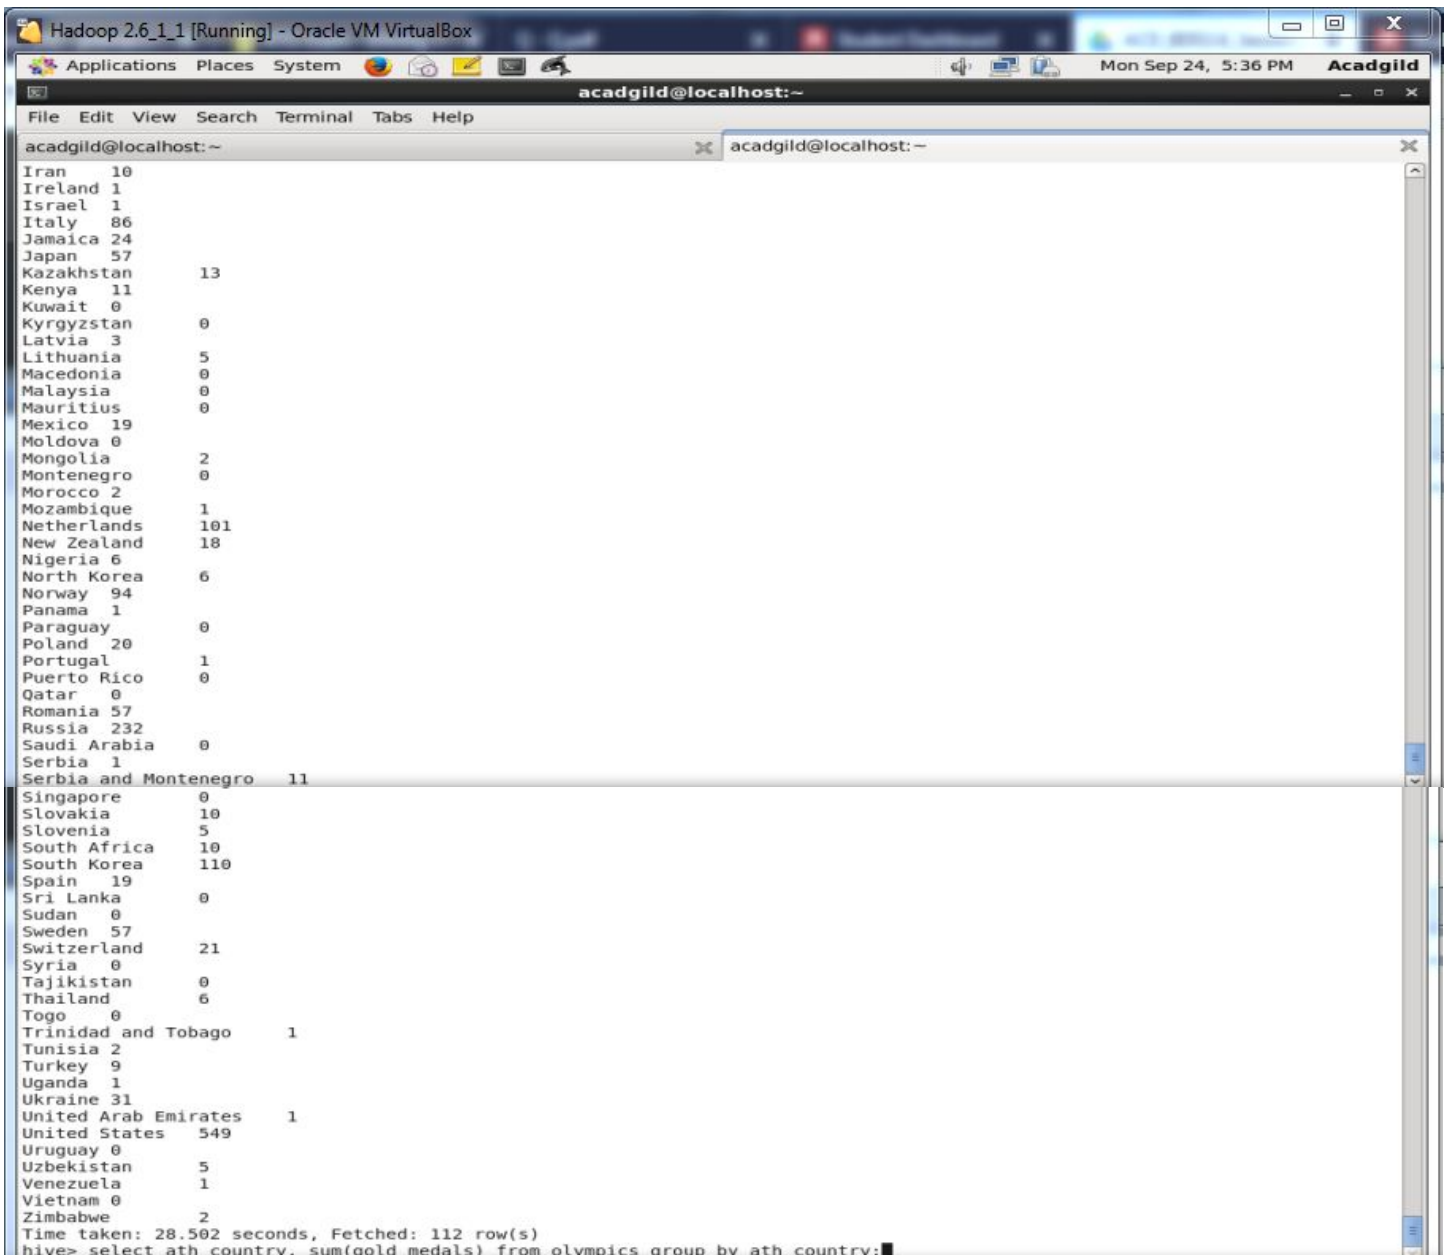
4. Write a hive program to find the number of gold medals each country own.

1. For this we need to write hive query as follows:

hive> select country, sum(gold_medals) from olympics group by ath_country;



```
acacgild@localhost:~  
hive> select country, sum(gold_medals) from olympics group by ath_country;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execu  
tion engine (i.e. spark, tez) or using Hive 1.X releases.  
Query ID = acacgild_20180924173246_ef15f761-ec2f-484a-bd09-f05be5963c94  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks not specified. Estimated from input data size: 1  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Starting Job = job_1537776060511_0007, Tracking URL = http://localhost:8088/proxy/application_1537776060511_0007/  
Kill Command = /home/acacgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537776060511_0007  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1  
2018-09-24 17:32:56,121 Stage-1 map = 0%, reduce = 0%  
2018-09-24 17:33:03,868 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.6 sec  
2018-09-24 17:33:12,529 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.02 sec  
MapReduce Total cumulative CPU time: 6 seconds 20 msec  
Ended Job = job_1537776060511_0007  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.02 sec HDFS Read: 528006 HDFS Write: 2741 SUCCESS  
Total MapReduce CPU Time Spent: 6 seconds 20 msec  
OK  
2004 2  
2008 2  
Afghanistan 0  
Algeria 2  
Argentina 49  
Armenia 0  
Australia 163  
Austria 36  
Azerbaijan 6  
Bahamas 11  
Bahrain 0  
Barbados 0  
Belarus 17  
Belgium 2  
Botswana 0  
Brazil 46  
Bulgaria 8  
Cameroon 20  
Canada 167  
Chile 3  
China 232  
Chinese Taipei 2  
Colombia 2  
Costa Rica 0  
Croatia 35  
Cuba 57  
Cyprus 0  
Czech Republic 14  
Denmark 46  
Dominican Republic 3  
Ecuador 0  
Egypt 1  
Eritrea 0  
Estonia 6  
Ethiopia 13  
Finland 11  
France 108  
Gabon 0  
Georgia 6  
Germany 223  
Great Britain 124  
Greece 12  
Grenada 1  
Guatemala 0  
Hong Kong 0  
Hungary 77  
Iceland 0  
India 1  
Indonesia 5
```

```
Hadoop 2.6.1_1 [Running] - Oracle VM VirtualBox
Applications Places System
acadmild@localhost:~
File Edit View Search Terminal Tabs Help
acadmild@localhost:~
Iran 10
Ireland 1
Israel 1
Italy 86
Jamaica 24
Japan 57
Kazakhstan 13
Kenya 11
Kuwait 0
Kyrgyzstan 0
Latvia 3
Lithuania 5
Macedonia 0
Malaysia 0
Mauritius 0
Mexico 19
Moldova 0
Mongolia 2
Montenegro 0
Morocco 2
Mozambique 1
Netherlands 101
New Zealand 18
Nigeria 6
North Korea 6
Norway 94
Panama 1
Paraguay 0
Poland 20
Portugal 1
Puerto Rico 0
Qatar 0
Romania 57
Russia 232
Saudi Arabia 0
Serbia 1
Serbia and Montenegro 11
Singapore 0
Slovakia 10
Slovenia 5
South Africa 10
South Korea 110
Spain 19
Sri Lanka 0
Sudan 0
Sweden 57
Switzerland 21
Syria 0
Tajikistan 0
Thailand 6
Togo 0
Trinidad and Tobago 1
Tunisia 2
Turkey 9
Uganda 1
Ukraine 31
United Arab Emirates 1
United States 549
Uruguay 0
Uzbekistan 5
Venezuela 1
Vietnam 0
Zimbabwe 2
Time taken: 28.502 seconds, Fetched: 112 row(s)
hive> select ath_country, sum(gold_medals) from olympics group by ath_country;
```

3. With this we can observe that the countries along with number of gold medals own.

Task 2 – Write a hive UDF that implements functionality of string concat_ws(string SEP, array<string>). This UDF will accept two arguments, one string and one array of string. It will return a single string where all the elements of the array are separated by the SEP.

1. For this task first we will write the UDF in Java and the source code is as follows:

```
import org.apache.hadoop.hive.ql.udf.generic;

import org.apache.hadoop.hive.ql.exec.Description;
import org.apache.hadoop.hive.ql.exec.UDFArgumentException;
import org.apache.hadoop.hive.ql.exec.UDFArgumentLengthException;
import org.apache.hadoop.hive.ql.exec.UDFArgumentTypeException;
import org.apache.hadoop.hive.ql.metadata.HiveException;
import org.apache.hadoop.hive.ql.udf.generic.GenericUDF;
import org.apache.hadoop.hive.ql.udf.generic.GenericUDF.DeferredObject;
```



```

import org.apache.hadoop.hive.serde.serdeConstants;
import org.apache.hadoop.hive.serde2.objectinspector.ListObjectInspector;
import org.apache.hadoop.hive.serde2.objectinspector.ObjectInspector;
import org.apache.hadoop.hive.serde2.objectinspector.ObjectInspector.Category;
import org.apache.hadoop.hive.serde2.objectinspector.PrimitiveObjectInspector;
import org.apache.hadoop.hive.serde2.objectinspector.PrimitiveObjectInspector.PrimitiveCategory;
import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorFactory;
import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorUtils;
import
org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorUtils.PrimitiveGrouping;
import org.apache.hadoop.io.Text;

```

```

public class GenericUDFConcatWS extends GenericUDF {
    private transient ObjectInspector[] argumentOIs;

```

```

    @Override

```

```

    public ObjectInspector initialize(ObjectInspector[] arguments) throws UDFArgumentException {
        if (arguments.length < 2) {
            throw new UDFArgumentLengthException(
                "The function CONCAT_WS(separator,[string | array(string)]+) "
                + "needs at least two arguments.");
        }
    }

```

```

    // check if argument is a string or an array of strings
    for (int i = 0; i < arguments.length; i++) {
        switch(arguments[i].getCategory()) {
            case LIST:
                if (isStringOrVoidType(
                    ((ListObjectInspector) arguments[i]).getListElementObjectInspector())) {
                    break;
                }
            case PRIMITIVE:
                if (isStringOrVoidType(arguments[i])) {
                    break;
                }
            default:
                throw new UDFArgumentTypeException(i, "Argument " + (i + 1)
                    + " of function CONCAT_WS must be \"" + serdeConstants.STRING_TYPE_NAME
                    + " or " + serdeConstants.LIST_TYPE_NAME + "<" +
serdeConstants.STRING_TYPE_NAME
                    + ">\", but \"" + arguments[i].getTypeName() + "\" was found.");
        }
    }

    argumentOIs = arguments;
    return PrimitiveObjectInspectorFactory.writableStringObjectInspector;
}

```

```

protected boolean isStringOrVoidType(ObjectInspector oi) {
    if (oi.getCategory() == Category.PRIMITIVE) {
        if (PrimitiveGrouping.STRING_GROUP

```

```

        == PrimitiveObjectInspectorUtils.getPrimitiveGrouping(
            ((PrimitiveObjectInspector) oi).getPrimitiveCategory())
        || ((PrimitiveObjectInspector) oi).getPrimitiveCategory() == PrimitiveCategory.VOID) {
    return true;
}
}
return false;
}

```

```

private final Text resultText = new Text();

```

```

@Override
public Object evaluate(DeferredObject[] arguments) throws HiveException {
    if (arguments[0].get() == null) {
        return null;
    }
    String separator = PrimitiveObjectInspectorUtils.getString(
        arguments[0].get(), (PrimitiveObjectInspector)argumentOIs[0]);

    StringBuilder sb = new StringBuilder();
    boolean first = true;
    for (int i = 1; i < arguments.length; i++) {
        if (arguments[i].get() != null) {
            if (first) {
                first = false;
            } else {
                sb.append(separator);
            }
            if (argumentOIs[i].getCategory().equals(Category.LIST)) {
                Object strArray = arguments[i].get();
                ListObjectInspector strArrayOI = (ListObjectInspector) argumentOIs[i];
                boolean strArrayFirst = true;
                for (int j = 0; j < strArrayOI.getListLength(strArray); j++) {
                    if (strArrayFirst) {
                        strArrayFirst = false;
                    } else {
                        sb.append(separator);
                    }
                    sb.append(strArrayOI.getListElement(strArray, j));
                }
            } else {
                sb.append(PrimitiveObjectInspectorUtils.getString(
                    arguments[i].get(), (PrimitiveObjectInspector)argumentOIs[i]));
            }
        }
    }

    resultText.set(sb.toString());
    return resultText;
}

```

```

@Override

```

```

public String getDisplayString(String[] children) {
    assert (children.length >= 2);
    return getStandardDisplayString("concat_ws", children);
}
}

```

- Now convert the above source code into *concatWS.jar* and load into hive as follows:

```
hive> jar /home/acadgild/concatWS.jar
```

The screenshot shows a terminal window titled "Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox" with the user "acadgild@localhost:~". The terminal displays the following commands and output:

```

hive> add jar /home/acadgild/concatWS.jar;
Added [/home/acadgild/concatWS.jar] to class path
Added resources: [/home/acadgild/concatWS.jar]
hive> list jars;
/home/acadgild/concatWS.jar

```

- And we can list the jars with *list jars* command.

- Now create a table *prema*, as follows:

The screenshot shows a terminal window titled "Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox" with the user "acadgild@localhost:~". The terminal displays the following commands and output:

```

hive> create table prema (
  > id int,
  > name string
  > )Row format delimited fields terminated by '\t';
OK
Time taken: 0.771 seconds

```

- Now the table is created and load some data into the table as follows:

The screenshot shows a terminal window titled "Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox" with the user "acadgild@localhost:~". The terminal displays the following commands and output:

```

hive> load data local inpath '/home/acadgild/prema.txt' into table prema;
Loading data to table sports.prema
OK
Time taken: 1.271 seconds

```

- Now select the contents from the table *prema* as follows:

```
hive>select * from prema;
```

The screenshot shows a terminal window titled "Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox" with the user "acadgild@localhost:~". The terminal displays the following commands and output:


```

hive> select * from prema;
OK
1      prema
2      vardhan
3      Reddy
Time taken: 0.297 seconds, Fetched: 3 row(s)

```

7. Now apply the *concat_ws* function on names which are present in the table *prema* as follows:

hive>select concat_ws(',', collect_list(name)) from prema;



```
acadgild@localhost:~  
hive> select concat_ws(',', collect_list(name)) from prema;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different ex  
tion engine (i.e. spark, tez) or using Hive 1.X releases.  
Query ID = acadgild_20180924231702_1a9a8595-98e8-469f-9344-c032bca7d899  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 1  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Starting Job = job_1537776060511_0011, Tracking URL = http://localhost:8088/proxy/application_1537776060511_0011/  
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537776060511_0011  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1  
2018-09-24 23:17:13,689 Stage-1 map = 0%, reduce = 0%  
2018-09-24 23:17:21,585 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.7 sec  
2018-09-24 23:17:31,250 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.05 sec  
MapReduce Total cumulative CPU time: 7 seconds 50 msec  
Ended Job = job_1537776060511_0011  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.05 sec HDFS Read: 8421 HDFS Write: 119 SUCCESS  
Total MapReduce CPU Time Spent: 7 seconds 50 msec  
OK  
prema, vardhan, Reddy  
Time taken: 29.832 seconds, Fetched: 1 row(s)  
hive>
```

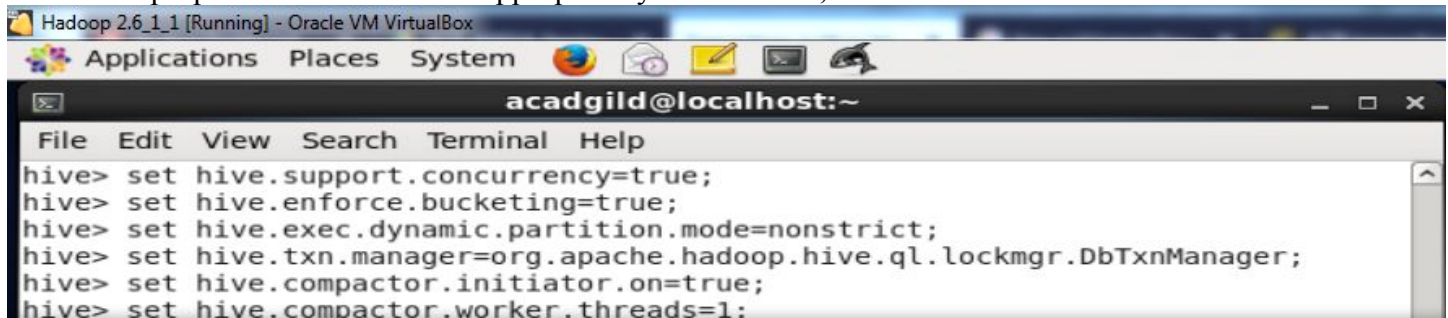
8. We can observe that in the output all the names in the table *prema* are separated by ‘,’ as *prema, vardhan, Reddy*

Task 3 – Row-level transactions available in Hive.

Transactions are provided at the row-level in Hive 0.14. The different row-level transactions available in Hive 0.14 are as follows:

1. Insert.
2. Delete.
3. Update.

The below properties needs to be set appropriately in *hive shell*, order-wise to work with transaction in Hive:



```
Hadoop 2.6.1.1 [Running] - Oracle VM VirtualBox  
Applications Places System  
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> set hive.support.concurrency=true;  
hive> set hive.enforce.bucketing=true;  
hive> set hive.exec.dynamic.partition.mode=nonstrict;  
hive> set hive.txn.manager=org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;  
hive> set hive.compactor.initiator.on=true;  
hive> set hive.compactor.worker.threads=1;
```

The sample file *sample.data*:

99001, Adam, California
99002, Brain, New York
99003, Crane, Chicago
99004, David, California
99005, Emily, New York
99006, Frank, Chicago
99007, George, Chicago
99008, Hall, New York
99009, Ivan, California
99010, Jacob, New York

Now create the table *emp_temp* that supports hive transactions as follows:

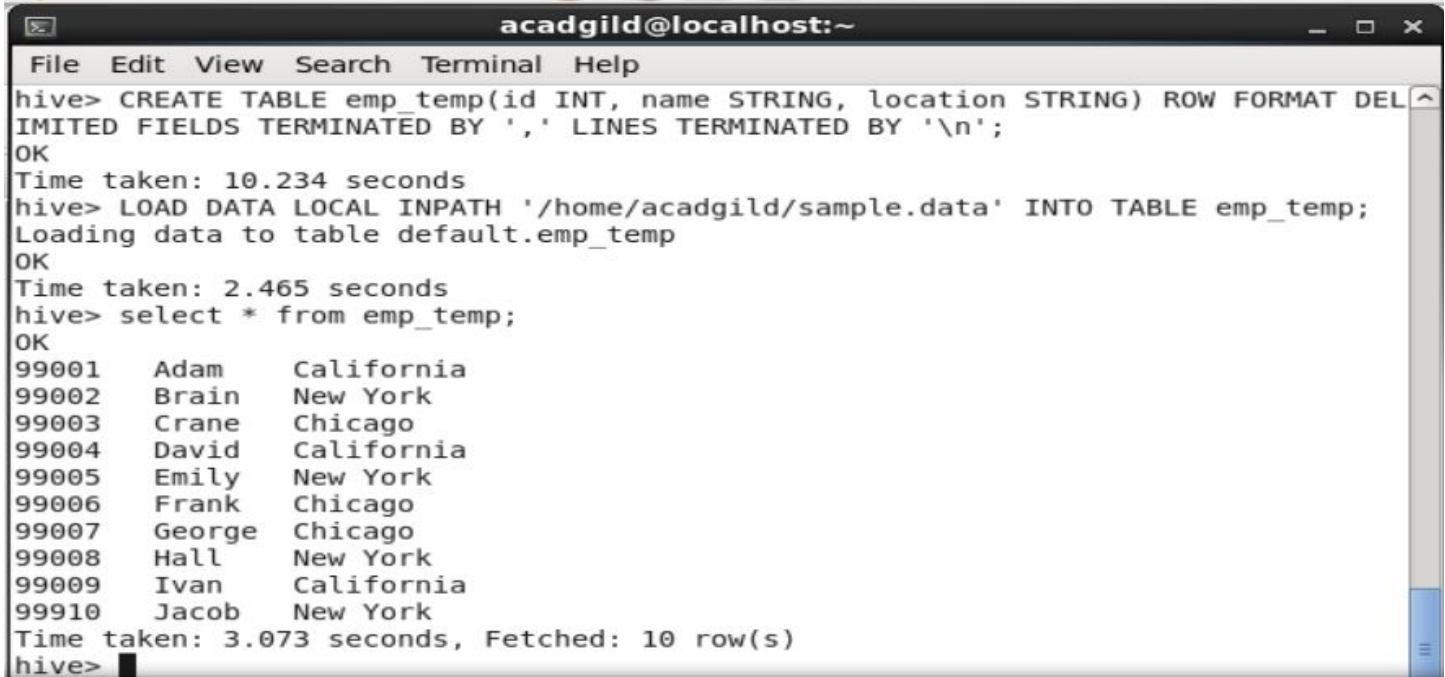
```
CREATE TABLE emp_temp(id INT, name STRING, location STRING)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
LINES TERMINATED BY '\n';
```

Load the data *sample.data* into the table *emp_temp* with the following command:

```
hive>LOAD DATA LOCAL INPATH '/home/acadgild/sample.data' INTO TABLE emp_temp;
```

Now display all the data that loaded into the table *emp_temp* as follows:

```
hive>select * from emp_temp;
```



```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> CREATE TABLE emp_temp(id INT, name STRING, location STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n';  
OK  
Time taken: 10.234 seconds  
hive> LOAD DATA LOCAL INPATH '/home/acadgild/sample.data' INTO TABLE emp_temp;  
Loading data to table default.emp_temp  
OK  
Time taken: 2.465 seconds  
hive> select * from emp_temp;  
OK  
99001 Adam California  
99002 Brain New York  
99003 Crane Chicago  
99004 David California  
99005 Emily New York  
99006 Frank Chicago  
99007 George Chicago  
99008 Hall New York  
99009 Ivan California  
99910 Jacob New York  
Time taken: 3.073 seconds, Fetched: 10 row(s)  
hive>
```

Now create another table called *emp_temp1* as follows:

```
hive>CREATE TABLE emp_temp1(id INT, name STRING, location STRING) CLUSTERED BY (ID) INTO 5  
BUCKETS ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n'  
STORED AS ORC TBLPROPERTIES('transactional'='true');
```



```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> CREATE TABLE emp_temp1(id INT, name STRING, location STRING) CLUSTERED BY (ID) INTO 5 BUCKETS ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n' STORED AS ORC TBLPROPERTIES('transactional'='true');  
OK  
Time taken: 0.23 seconds
```

Now load data into table '*emp_temp1*' from table '*emp_temp*': as follows:

```
hive> FROM emp_temp INSERT INTO emp_temp1 SELECT id,name,location ORDER BY 1;
```

```
acadgild@localhost:~
File Edit View Search Terminal Help
hive> FROM emp_temp INSERT INTO emp_temp1 SELECT id, name, location ORDER BY 1;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180925121119_d63a3347-2064-4d1d-a172-d4b8a08c7cb6
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537851853748_0001, Tracking URL = http://localhost:8088/proxy/application_1537851853748_0001/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537851853748_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-09-25 12:11:37,256 Stage-1 map = 0%, reduce = 0%
2018-09-25 12:11:45,353 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.47 sec
2018-09-25 12:11:54,530 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.37 sec
ec
2018-09-25 12:12:42,684 Stage-2 map = 100%, reduce = 20%, Cumulative CPU 5.62 sec
ec
2018-09-25 12:12:43,990 Stage-2 map = 100%, reduce = 33%, Cumulative CPU 7.86 sec
ec
2018-09-25 12:12:48,032 Stage-2 map = 100%, reduce = 53%, Cumulative CPU 11.97 sec
2018-09-25 12:12:49,199 Stage-2 map = 100%, reduce = 67%, Cumulative CPU 14.12 sec
2018-09-25 12:12:50,405 Stage-2 map = 100%, reduce = 87%, Cumulative CPU 17.75 sec
2018-09-25 12:12:51,521 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 20.56 sec
MapReduce Total cumulative CPU time: 20 seconds 560 msec
Ended Job = job_1537851853748_0002
Loading data to table default.emp_temp1
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.37 sec HDFS Read: 6744 HDFS Write: 457 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 5 Cumulative CPU: 20.56 sec HDFS Read: 19955 HDFS Write: 4539 SUCCESS
Total MapReduce CPU Time Spent: 25 seconds 930 msec
OK
Time taken: 93.754 seconds
```

Now print the elements from the table as follows:

hive>select *from emp_temp1;

```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> select * from emp_temp1;  
OK  
99910      Jacob      New York  
99005      Emily      New York  
99006      Frank      Chicago  
99001      Adam      California  
99007      George     Chicago  
99002      Brain     New York  
99008      Hall      New York  
99003      Crane     Chicago  
99009      Ivan      California  
99004      David     California  
Time taken: 0.457 seconds, Fetched: 10 row(s)  
hive> █
```

Now if we try to insert the same data again, it will be append to the previous data as shown below:

```
hive> FROM emp_temp INSERT INTO emp_temp1 SELECT id, name, location ORDER BY 1;  
hive> FROM emp_temp INSERT INTO emp_temp1 SELECT id, name, location ORDER BY 1;  
  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.  
Query ID = acadgild_20180925123636_b3ff2c34-869d-4451-85e3-b544783f1fda  
Total jobs = 2  
Launching Job 1 out of 2  
Number of reduce tasks determined at compile time: 1  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Starting Job = job_1537851853748_0003, Tracking URL = http://localhost:8088/proxy/application_1537851853748_0003/  
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537851853748_0003  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1  
2018-09-25 12:36:47,794 Stage-1 map = 0%, reduce = 0%  
2018-09-25 12:36:55,692 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.14 sec  
2018-09-25 12:37:03,352 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.1 sec  
MapReduce Total cumulative CPU time: 5 seconds 100 msec  
Ended Job = job_1537851853748_0003  
Launching Job 2 out of 2  
Number of reduce tasks determined at compile time: 5  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:
```



```

y/application_1537851853748_0004/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill
job_1537851853748_0004
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 5
2018-09-25 12:37:19,248 Stage-2 map = 0%, reduce = 0%
2018-09-25 12:37:27,021 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 1.68 se
c
2018-09-25 12:37:45,592 Stage-2 map = 100%, reduce = 13%, Cumulative CPU 3.8 se
c
2018-09-25 12:37:50,650 Stage-2 map = 100%, reduce = 20%, Cumulative CPU 5.49 s
ec
2018-09-25 12:37:51,866 Stage-2 map = 100%, reduce = 33%, Cumulative CPU 7.54 s
ec
2018-09-25 12:37:54,310 Stage-2 map = 100%, reduce = 40%, Cumulative CPU 9.11 s
ec
2018-09-25 12:37:55,522 Stage-2 map = 100%, reduce = 67%, Cumulative CPU 13.82
sec
2018-09-25 12:37:56,756 Stage-2 map = 100%, reduce = 83%, Cumulative CPU 16.51
sec
2018-09-25 12:37:57,844 Stage-2 map = 100%, reduce = 97%, Cumulative CPU 19.57
sec
2018-09-25 12:37:58,924 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 20.75
sec
MapReduce Total cumulative CPU time: 20 seconds 750 msec
Ended Job = job_1537851853748_0004
Loading data to table default.emp_temp1
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.1 sec HDFS Read: 6744 HDF
S Write: 457 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 5 Cumulative CPU: 20.75 sec HDFS Read: 19740
HDFS Write: 4546 SUCCESS
Total MapReduce CPU Time Spent: 25 seconds 850 msec
OK
Time taken: 83.983 seconds

```

Now use print the data from table as follows:

```

hive> select * from emp_temp1;
hive> select * from emp_temp1;
OK
99910    Jacob    New York
99005    Emily    New York
99910    Jacob    New York
99005    Emily    New York
99006    Frank    Chicago
99001    Adam     California
99006    Frank    Chicago
99001    Adam     California
99007    George   Chicago
99002    Brain    New York
99007    George   Chicago
99002    Brain    New York
99008    Hall     New York
99003    Crane    Chicago
99008    Hall     New York
99003    Crane    Chicago
99009    Ivan     California
99004    David    California
99009    Ivan     California
99004    David    California
Time taken: 0.29 seconds, Fetched: 20 row(s)
hive>

```

Updating the Data in Hive Table

Update the table record as follows:

```

hive> UPDATE emp_temp1 SET location='Delhi' WHERE location='New York';

```



```

hive> UPDATE emp_temp1 SET location='Delhi' WHERE location='New York';
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180925124604_5e22cdb2-1d45-4063-818c-cb8dcc78f8ef
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537851853748_0005, Tracking URL = http://localhost:8088/proxy/application_1537851853748_0005/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537851853748_0005
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-09-25 12:46:14,650 Stage-1 map = 0%, reduce = 0%
2018-09-25 12:46:41,960 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 4.51 sec
2018-09-25 12:46:45,683 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 13.11 sec
2018-09-25 12:46:48,263 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 18.74 sec
2018-09-25 12:46:49,504 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 23.66 sec
2018-09-25 12:46:51,977 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 24.48 sec
2018-09-25 12:47:11,315 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 26.4 sec
2018-09-25 12:47:12,601 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 28.53 sec
2018-09-25 12:47:13,856 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 30.65 sec
2018-09-25 12:46:41,960 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 4.51 sec
2018-09-25 12:46:45,683 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 13.11 sec
2018-09-25 12:46:48,263 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 18.74 sec
2018-09-25 12:46:49,504 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 23.66 sec
2018-09-25 12:46:51,977 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 24.48 sec
2018-09-25 12:47:11,315 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 26.4 sec
2018-09-25 12:47:12,601 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 28.53 sec
2018-09-25 12:47:13,856 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 30.65 sec
2018-09-25 12:47:15,133 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 32.8 sec
2018-09-25 12:47:16,417 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 34.82 sec
2018-09-25 12:47:17,689 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 38.68 sec
2018-09-25 12:47:18,904 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 42.11 sec
2018-09-25 12:47:20,035 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 43.89 sec
MapReduce Total cumulative CPU time: 43 seconds 890 msec
Ended Job = job_1537851853748_0005
Loading data to table default.emp_temp1
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 43.89 sec HDFS Read: 58267
HDFS Write: 2699 SUCCESS
Total MapReduce CPU Time Spent: 43 seconds 890 msec
OK
Time taken: 79.003 seconds

```

This will update the Records, with 'New York' to 'Delhi' we will see the output as follows:

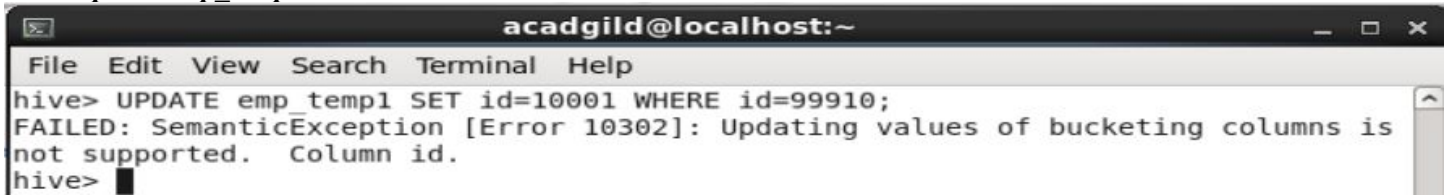
```
hive>select * from emp_temp1;
```

```
hive> SELECT * FROM EMP_TEMP1;
OK
99910   Jacob   Delhi
99005   Emily   Delhi
99910   Jacob   Delhi
99005   Emily   Delhi
99006   Frank   Chicago
99001   Adam    California
99006   Frank   Chicago
99001   Adam    California
99007   George  Chicago
99002   Brain   Delhi
99007   George  Chicago
99002   Brain   Delhi
99008   Hall    Delhi
99003   Crane   Chicago
99008   Hall    Delhi
99003   Crane   Chicago
99009   Ivan    California
99004   David   California
99009   Ivan    California
99004   David   California
Time taken: 0.262 seconds, Fetched: 20 row(s)
hive>
```

We have successfully updated the data, and we can observe that the updated data.

Now we try to update the ID of the employee as follows:

hive> update emp_temp1 set id=10001 where id=99910.



The screenshot shows a terminal window titled 'acadgild@localhost:~'. The terminal contains the following text:

```
File Edit View Search Terminal Help
hive> UPDATE emp_temp1 SET id=10001 WHERE id=99910;
FAILED: SemanticException [Error 10302]: Updating values of bucketing columns is
not supported. Column id.
hive>
```

From the above image we can see that we have received an error message. This means that the Update command is not supported on the columns that are bucketed.

Deleting a Row from Hive Table

Now let's perform the Delete operation on the same table as follows:

hive>DELETE FROM emp_temp1 WHERE id=99910;

```

hive> DELETE FROM emp_temp1 WHERE id=99910;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180925130215_f159a1e7-04cf-4cbb-8eb3-645f78994496
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1537851853748_0006, Tracking URL = http://localhost:8088/proxy/application_1537851853748_0006/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1537851853748_0006
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-09-25 13:02:25,010 Stage-1 map = 0%, reduce = 0%
2018-09-25 13:02:50,586 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 7.54 sec
2018-09-25 13:02:54,260 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 8.91 sec
2018-09-25 13:02:58,245 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 19.84 sec
2018-09-25 13:03:00,885 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 21.04 sec
2018-09-25 13:03:02,112 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 22.42 sec
2018-09-25 13:03:19,917 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 24.6 sec
2018-09-25 13:03:21,202 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 26.68 sec
2018-09-25 13:03:22,516 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 28.84 sec
2018-09-25 13:03:25,106 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 34.22 sec
2018-09-25 13:03:26,284 Stage-1 map = 100%, reduce = 87%, Cumulative CPU 37.84 sec
2018-09-25 13:03:27,379 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 40.63 sec
MapReduce Total cumulative CPU time: 40 seconds 630 msec
Ended Job = job_1537851853748_0006
Loading data to table default.emp_temp1
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 40.63 sec HDFS Read: 57442 HDFS Write: 765 SUCCESS
Total MapReduce CPU Time Spent: 40 seconds 630 msec
OK
Time taken: 74.113 seconds
hive>

```

Now we can see that the record was deleted. We can observe this by printing the table records as follows:

```
hive> select * from emp_temp1;
```

```
hive> select * from emp_temp1;
OK
99005    Emily    Delhi
99005    Emily    Delhi
99006    Frank    Chicago
99001    Adam     California
99006    Frank    Chicago
99001    Adam     California
99007    George   Chicago
99002    Brain    Delhi
99007    George   Chicago
99002    Brain    Delhi
99008    Hall     Delhi
99003    Crane    Chicago
99008    Hall     Delhi
99003    Crane    Chicago
99009    Ivan     California
99004    David    California
99009    Ivan     California
99004    David    California
Time taken: 0.228 seconds, Fetched: 18 row(s)
hive>
```

We can see that there is no row with *id=99910*. This means that we have successfully deleted the row from the Hive table.

This is how the transactions or row-wise operations are performed in Hive.